

# CHAPTER 1

---

## Graphs and gating functions

---

*J. Stanier*

Progress: 90%

Text and figures forming in progress

.....

### 1.1 Introduction

Many compilers represent the input program as some form of graph in order to aid analysis and transformation. A cornucopia of program graphs have been presented in the literature and implemented in real compilers. Therefore it comes as no surprise that a number of program graphs use SSA concepts as the core principle of their representation. These range from very literal translations of SSA into graph form to more abstract graphs which are implicitly SSA. We aim to introduce a selection of program graphs which use SSA concepts, and examine how they may be useful to a compiler writer.

One of the seminal graph representations is the Control Flow Graph (CFG), which was introduced by Allen to explicitly represent possible control paths in a program. Traditionally, the CFG is used to convert a program into SSA form. Additionally, representing a program in this way makes a number of operations simpler to perform, such as identifying loops, discovering irreducibility and performing interval analysis techniques.

The CFG models control flow, but many graphs model *data flow*. This is useful as a large number of compiler optimizations are based on data flow. The graphs we consider in this section are all data flow graphs, representing the data dependencies in a program. We will look at a number of different SSA-based graph representations. These range from those which are a very literal translation of SSA into a graph form to those which are more abstract in nature. An introduction to each graph will be given, along with diagrams to show how sample programs look

when translated into that particular graph. Additionally, we will touch on the literature describing a usage of a given graph with the application that it was used for.

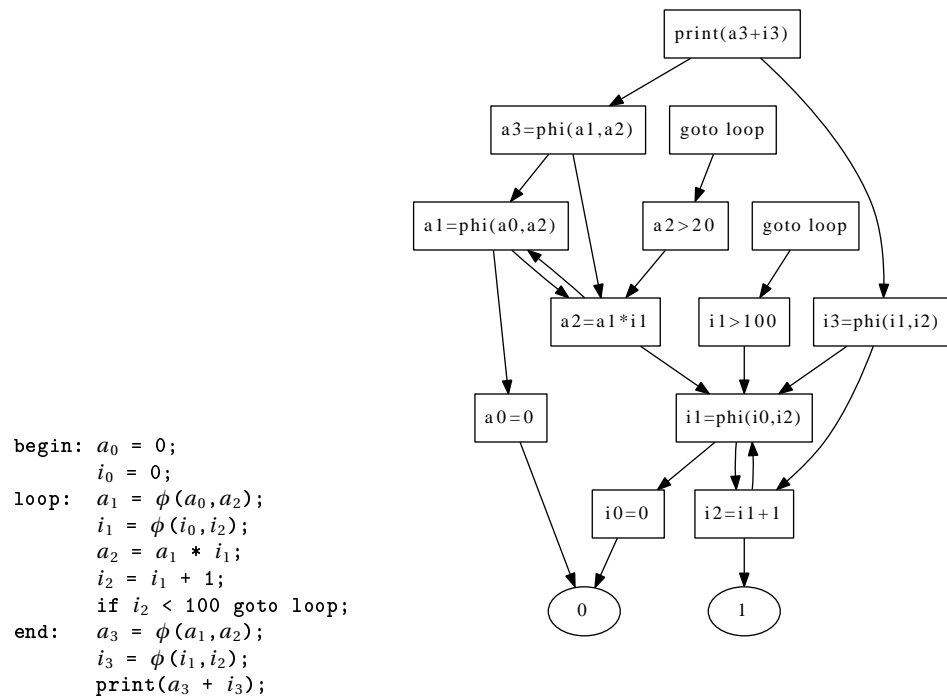
## 1.2 Data Flow Graphs

The data flow graph (DFG) is also a directed graph  $G = (V, E)$  except edges  $E$  represent the flow of data from the result of one operation to the input of another. An instruction executes once all of its input data values have been consumed. When an instruction executes it produces a new data value which is propagated to other connected instructions.

Whereas the CFG imposes a total ordering on instructions, the DFG has no such concept, nor does the DFG contain whole program information. Thus, target code cannot be generated directly from the DFG. The DFG can be seen as a companion to the CFG, and they can be generated alongside each other. With access to both graphs, optimisations such as dead code elimination, constant folding and common subexpression elimination can be performed effectively. However, keeping both graphs updated during optimisation can be costly and complicated.

## 1.3 The SSA Graph

We begin our exploration with a graph that is very similar to SSA: the SSA Graph. Many different variations exist in the literature, so we take ours from Cooper et al. An SSA Graph consists of vertices which represent operations (such as add and load) or  $\phi$ -functions, and directed edges connect uses to definitions of values. The edges to a vertex represent the arguments required for that operation, and the edge from a vertex represents the propagation of that operation's result after it has been computed. This graph is therefore a *demand-based* representation. In order to compute a vertex, we must first *demand* the results of the operands and then perform the operation indicated on that vertex. Reversing the direction of each edge would provide the widely used *data-based* representation. The SSA Graph can be constructed from a program in SSA form by *explicitly* adding use-definition chains. We present some sample code in Figure 1.1 which is then translated into an SSA Graph. ~~Note that there are no explicit nodes for variables in the graph. Instead, an operator node can be seen as the “location” of the value stored in a variable. We have annotated operators with variable names to show the correspondence between the operations in the graph and in the SSA form program.~~ Note that each node mix up the operation and the variable(s) it defines, as actual data structures might enable to find one from the other.



**Fig. 1.1** Some SSA code translated into an SSA graph.

The textual representation of SSA is much easier for a human to read. However, the primary benefit of representing the input program in this form is that the compiler writer is able to apply a wide array of graph-based optimizations by using standard graph traversal and transformation techniques. It is possible to augment the SSA Graph to model memory dependencies. This ~~is~~ can be achieved by adding additional *state edges* that enforce an order of interpretation. These edges are extensively used in the Value State Dependence Graph, which we will look at later, after we touch on the concept of gating functions.

In the literature, the SSA Graph has been used to detect a variety of induction variables in loops, also for performing instruction selection techniques, operator strength reduction, rematerialization, and has been combined with an extended SSA language to aid compilation in a parallelizing compiler. The reader should note that the exact specification of what constitutes an SSA Graph changes from paper to paper. The essence of the IR has been presented here, as each author tends to make small modifications for their particular implementation.

### 1.3.1 Finding induction variables with the SSA Graph

We illustrate the usefulness of the SSA Graph through a basic induction variable recognition technique. A more sophisticated technique is developed in Chapter ??. Given that a program is represented as an SSA Graph, the task of finding induction variables is simplified. A *basic linear induction variable*  $i$  is a variable that appears only in the form:

```

i = 10
loop
    ...
    i = i + k
    ...
endloop

```

i is a basic induction variable.

where  $k$  is a constant or loop invariant. Wolfe made the observation that each basic linear induction variable will belong to a ~~nontrivial strongly connected region (SCR)~~ non-trivial strongly connected component (SCC) in the SSA graph. ~~There exist linear algorithms for discovering SCRs~~ SCCs can be easily discovered in linear time using any depth first search traversal. Each such ~~SCR~~ SCC must conform to the following constraints:

- The ~~SCR~~ SCC contains only one  $\phi$ -function at the header of the loop.
- The ~~SCR~~ SCC contains only addition and subtraction operators, and the right operand of the subtraction is not part of the ~~SCR~~ SCC (no  $i = n - i$  assignments).
- The other operand of each addition or subtraction is loop invariant.

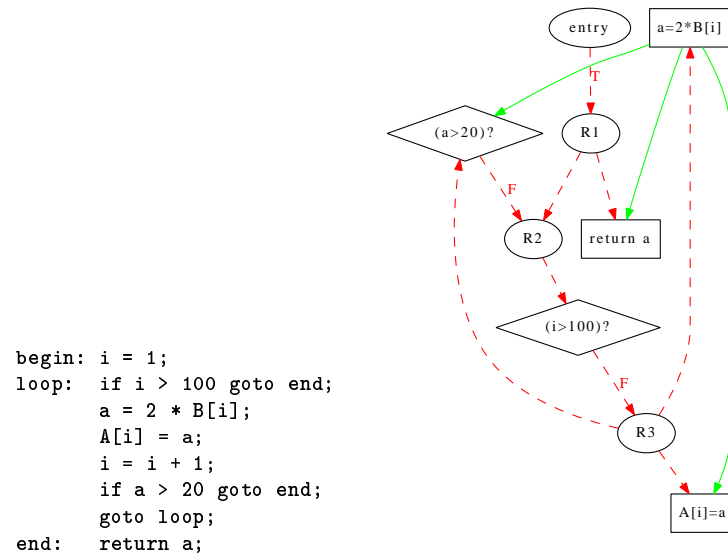
This technique can be expanded to detect a variety of other classes of induction variables, such as wrap-around variables, non-linear induction variables and nested induction variables. Scans and reductions also show a similar SSA Graph pattern and can be detected using the same approach.

## 1.4 Program Dependence Graph

The Program Dependence Graph (PDG) represents both control and data dependencies together in one graph. The PDG was developed to aid ~~optimisations~~ optimizations requiring reordering of instructions and graph rewriting for parallelism, as the strict ordering of the CFG is relaxed and complemented by the presence of data dependence information. The PDG is a directed graph  $G = (V, E)$  where nodes  $V$  are statements, predicate expressions or region nodes, and edges  $E$  represent either control or data dependencies. Thus, the set of all edges  $E$  has two distinct subsets: the control dependence subgraph  $E_C$  and the data dependence subgraph  $E_D$ .  ~~$E_C$  can be cyclic if a loop is present in the program, since a loop in the PDG is defined by a control back edge forming an SCR.  $E_D$  is always acyclic, and can be seen as a series of data dependency DAGs for each basic block, which are then connected together based on the data flow through the program.~~ Similar to the CFG, a PDG also has two nodes ENTRY and EXIT, through which control flow enters and exits the program respectively. It is assumed that every node of the CFG is reachable from the entry node and can reach the exit node.

Statement nodes represent instructions in the program. Predicate nodes test a conditional statement and have true and false edges to represent the choice taken on evaluation of the predicate. Region nodes group ~~all nodes with the same control dependencies together, and order them into a hierarchy~~ control dependencies with identical source and label together. If the control dependence for a region node is satisfied, then it follows that all of its children can be executed. Thus, if a region node has three different control-independent statements as immediate children, then these could potentially be executed in parallel. Diagrammatically, rectangular nodes represent statements, diamond nodes predicates, and circular nodes are region nodes. ~~Filled-Dashed~~ edges represent control dependence, and ~~dashed-solid~~ edges represent data dependence. Loops in the PDG are represented by back edges in the control dependence subgraph. We show example code translated into a PDG in Figure 1.2.

Construction of the PDG is tackled in two steps from the CFG: construction of the control dependence subgraph and construction of the data dependence subgraph. ~~Ferrante et al. construct~~ The construction of the control dependence subgraph in  $O(N^2)$  time falls itself into two steps. First, control dependence between statements and predicate nodes is computed; then region nodes are added. A node  $w$  is said to be control dependent on node  $u$  along CFG edge  $(u, v)$  or simply control dependent on edge  $(u, v)$  if  $w$  post-dominates  $v$  and  $w$  does not strictly post-dominate  $u$ . Control dependencies between nodes is nothing else than post-dominance frontier i.e. dominance frontier on the reverse CFG. A slight difference in its construction process is that the corresponding control dependence edges from  $u$  to  $w$  is labelled by the boolean value taken by the predicate computed in  $u$  when branching on edge  $(u, v)$ . To compute the control dependence subgraph, the postdominator tree is constructed for the procedure. Then, the ENTRY node is added with one edge labelled  ~~$T$~~  true pointing to the CFG



**Fig. 1.2** Some code translated into a PDG. Nodes associated with the evolution of the induction variable  $i$  are omitted.

entry node, and another labelled ~~F~~-false going to the CFG exit node. Then, let  $S$  consist of all edges  $(A, B)$  in the CFG such that  $B$  is not an ancestor of  $A$  in the postdominator tree. Each of these edges has an associated label ~~T~~-or-~~F~~true or false. Then, each edge in  $S$  is considered in turn. Given  $(A, B)$ , the postdominator tree is traversed backwards from  $B$  until we reach  $A$ 's parent, marking all nodes visited (including  $B$ ) as control dependent on  $A$  with the label of  $S$ .

Next, region nodes are added to the PDG. Each region node summarizes ~~the~~ a set of control conditions and ~~groups~~-groups all nodes with the same set of control conditions together. Region nodes are also inserted so that predicate nodes will only have two successors, ~~thus hierarchically ordering the control dependencies~~. To begin with, an unpruned PDG is created by checking, for each node of the CFG, which control region ~~or control dependence node it belongs to it depends on~~. This is done by traversing the postdominator tree in postorder,

and using a hash table to map sets of control dependencies to region nodes. For each node  $N$  visited in the postdominator tree, the hash table is checked for an existing region node with the same set  $CD$  of control dependencies. If none exists, a new region node  $R$  is created with these control dependencies and entered into the hash table.  $R$  is made to be the only control dependence predecessor of  $N$ . Next, the intersection  $INT$  of  $CD$  is computed for each immediate child of  $N$  in the postdominator tree. If  $INT = CD$  then the corresponding dependencies are deleted from the child and replaced with a single dependence on the child's control predecessor. Then, a pass over the graph is made to make sure that each predicate node has a unique successor for each truth value. If more than one exists, ~~a region node is created and inserted as the successor of that particular predicate node's  $T$  or  $F$  edge~~ the corresponding edges are replaced by a single edge to a freshly created region node that itself points to the successor nodes.

~~The data dependence subgraph can be constructed after aliasing, procedure calls and side effects are analysed in the program. DAGs are constructed for each basic block during the initial parse of the source program. Each upwards exposed use in a block has a corresponding DAG leaf node as usual; these leaves are called merge nodes. Data flow analysis is then performed to compute the set of reaching definitions for each basic block, with the additional assumption that every variable is initially assigned the value "undefined" at program entry. Finally, the individual DAGs are connected to one another using the results of the data flow computation: edges are added from definition nodes to the corresponding merge nodes that may be reached. A statement  $B$  that is to be executed after a statement  $A$  in the original sequential ordering of the program depends on  $A$ , in the following situations: (flow)  $B$  reads to a storage location that was lastly accessed by  $A$  through a write; (anti)  $B$  writes to a storage location previously accessed through a read by  $A$ ; (output)  $A$  and  $B$  both have a write access to the same storage location. Side effects can also dictate the insertion of a dependence between  $A$  and  $B$  to force the sequential ordering of the final schedule. Memory accesses can not always be analyzed with enough precision. In the presence of a may-alias between two consecutive accesses a conservative dependence is to be inserted also. Memory access locations often vary with the iterations of the enclosing loops. Dependence analysis can take advantage of some abstract representations of the access function such as when the memory address can be represented as an affine function of the induction variables. Not only this enables a refinement of the alias information, but also the dependence can be labelled with a distance vector as a function itself of the loop indices. As an example, a loop indexed by  $i$ , that would access array  $A[i]$  twice, first as a write at iteration  $i$ , then as a read at iteration  $2i + 1$ , would lead to a flow dependence of distance  $i + 1$ .~~

~~Harrold et al. construct the PDG~~ The PDG can also be constructed during parsing. ~~Ferrante et al. considered the problem of generating linear code~~ Sequential code can be derived from a PDG, showing that but generating the *minimal size* CFG from a PDG ~~is turns out to be~~ an NP-Complete problem. ~~Further work has improved on the efficiency and generality of generating sequential code.~~ The

PDG's structure has been exploited for generating code for vectorisation, and has also been used in order to perform accurate program slicing and testing.

### 1.4.1 Detecting parallelism with the PDG

The structure of the PDG allows for parallelism to be detected easily. ~~In order to discover potential parallelism from the CFG, the~~ On a regular CFG representation, a scheduler based on data dependencies will generally restrict its scope to hyper-blocks. ~~In this context, code transformations such as loop unrolling or if-conversion transformation must first be performed (see Chapter ??) that effectively change control dependencies into data dependencies can expose instruction level parallelism.~~ However, the PDG can *directly* be used to detect parallelism. ~~Any node in the PDG~~ As an example, ~~any node of a CFG loop~~ that is not contained in an ~~SCR-SCC of the PDG (consisting of both control and data dependencies can be parallelized) can be vectorized.~~ In the example in Figure ?? ~~we have a program containing a loop consisting of multiple assignment statements, and the corresponding PDG. Here~~ 1,2, since the ~~statements~~ **statement  $d = e + f$  and  $g = h + i$**  (outlined in bold) in the loop do not form an ~~SCR~~, so they are parallelizable. The statements  ~~$a = b + c$  and  $i = i + a$~~  are not, since the assignment to  $i$  is data dependent on the assignment to ~~SCC of the PDG, it can be vectorized provided array expansion of variable  $a$ . Because of the circuit involving the test on  $a$ , the statement  $a = 2 * B[i]$  cannot.~~

Detecting parallelism with the PDG. The statements  ~~$d = e + f$  and  $g = h + i$~~  (outlined in bold) are parallelizable.

.....

## 1.5 Gating functions and GSA

In SSA form,  $\phi$ -functions are used to identify points where variable definitions converge. However, they cannot be directly *interpreted*, as they do not specify the condition which determines which of the variable definitions to choose. By this logic, we cannot directly interpret the SSA Graph. Being able to interpret our IR is a useful property as it gives the compiler writer more information when implementing optimizations, and also reduces the complexity of performing code generation. Gated Single Assignment form (GSA; sometimes called Gated SSA) is an extension of SSA with *gating functions*. These gating functions are directly interpretable versions of  $\phi$ -nodes, and replace  $\phi$ -nodes in the representation. ~~There are three~~ We usually distinguish the three following forms of gating function and we take our definition from Tu and Padua functions:

- The  ~~$\gamma$~~  $\phi_{if}$  function explicitly represents the condition which determines which  $\phi$  value to select. A  ~~$\gamma$~~ function is  ~~$\phi_{if}$~~ function of the form  ~~$\gamma(P, V_1, V_2)$~~  $\phi_{if}(P, V_1, V_2)$



where  $P$  is a predicate, and  $V_1$  and  $V_2$  are the values to be selected if the predicate evaluates to true or false respectively. This can be read simply as *if-then-else*.

- The  $\mu\phi_{entry}$  function is inserted at loop headers to select the initial and loop carried values. A  $\mu\phi_{entry}$  function is of the form  $\mu(V_{init}, V_{iter})\phi_{entry}(V_{init}, V_{iter})$ , where  $V_{init}$  is the initial input value for the loop, and  $V_{iter}$  is the iterative input. We replace  $\phi$ -functions at loop headers with  $\mu\phi_{entry}$  functions.
- The  $\eta\phi_{exit}$  function determines the value of a variable when a loop terminates. An  $\eta\phi_{exit}$  function is of the form  $\eta(P, V_{final})\phi_{exit}(P, V_{final})$  where  $P$  is a predicate and  $V_{final}$  is the definition reaching beyond the loop.

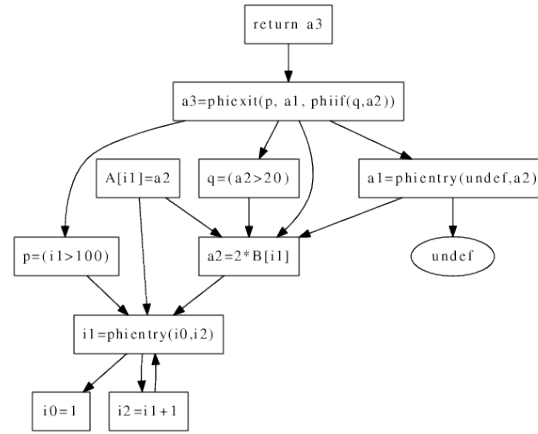


Fig. 1.3 A graph representation of our sample code in GSA form.

It is easiest to understand these gating functions by means of an example. Figure 1.3 shows how our earlier code in Figure 1.1–1.2 translates into GSA form. Here, we can see the use of both  $\mu$  and  $\eta\phi_{entry}$  and  $\phi_{exit}$  gating functions. At the header of our sample loop the  $\phi$ -functions have been replaced by  $\mu$  functions a  $\phi_{entry}$  function which determine between the initial and iterative values of  $a$  and  $i$  value of  $i$ . After the loop has finished executing, the two  $\eta$  functions propagate the correct value from the corresponding  $\mu$  function.

These gating functions are important as the concept will form components of the Value State Dependence Graph later. GSA has seen a number of uses in the literature. The original usage of GSA was by Ballance et al. as an intermediate stage in the construction of the Program Dependence Web IR. Havlak presented an algorithm for construction of a simpler version of GSA—Thinned GSA—which is constructed from a CFG in SSA form. GSA has been used for a number of analyses and transformations based on data flow nested  $\phi_{exit}$  functions selects the correct live-out version of  $a$ .

Tu and Padua presented an algorithm that constructs SSA and GSA simultaneously in a single process. Construction of GSA proceeds from the CFG. It uses the

concept of *path expressions*. A path expression  $P$  of type  $(u, v)$  is a simple regular expression over  $E$  such that every string in  $\sigma(P)$  is a path from node  $u$  to node  $v$  (where  $\sigma(P)$  represents the string generated by the regular expression  $P$ ). These path expressions are represented as *gating functions*. This example shows several interesting points. First, the semantic of both the  $\phi_{exit}$  and  $\phi_{if}$  are strict in their gate: here  $a_1$  or  $\phi_{exit}(q, a_2)$  are not evaluated before  $p$  is known<sup>1</sup>. Similarly, a  $\phi_{if}$  function that results from the nested if-then-else code of Figure 1.4 would be itself nested as  $a = \phi_{if}(p, \phi_{if}(q, a_2, a_3), a_1)$ . Second, this representation of the program does not allow for an interpreter to decide whether an instruction with a side effect (such as  $A[i_1] = a_2$  in our running example) has to be executed or not. Finally, computing the values of gates is highly related to the simplification of path expressions: in our running example  $a_2$  should be selected when the path  $\neg p$  followed by  $q$  (denoted  $\neg p.q$ ) is taken while  $a_1$  should be selected when the path  $p$  is taken; for our nested if-then-else example,  $a_1$  should be selected either when the path  $\neg p.r$  is taken or when the path  $\neg p.\neg r$  is taken which simplifies to  $\neg p$ . Diverse approaches can be used to generate the correct nested  $\phi_{if}$  or  $\phi_{exit}$  gating functions.

The symbol  $A$  represents an unconditional edge. The symbol  $\emptyset$  represents an edge not taken at a branch node. A  $\gamma$  expression  $\gamma(P, e_1, e_2, \dots, e_n)$  where only one  $e_i$  is  $A$  and all the other  $e$ 's are  $\emptyset$  represents the  $i$ 's edge from an  $n$ -way branch statement with condition  $P$ .

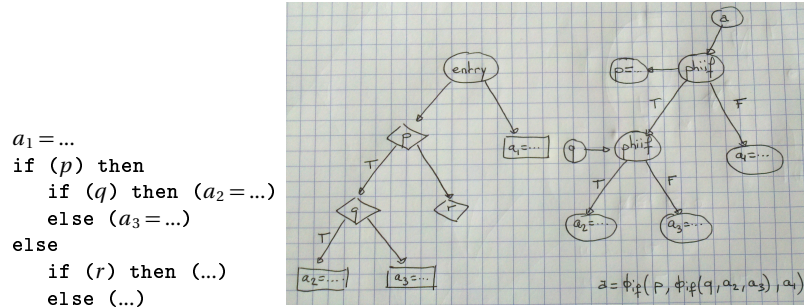


Fig. 1.4 (a) A structured code; (b) the CDG (with region nodes omitted); (c) the DAG representation of the nested gated  $\phi_{if}$ .

Then, these gating functions can be simplified along gating paths. The full algorithm for construction and simplification of gating functions can be found in the Tu and Padua paper in The more natural way uses a data flow analysis that, for each program points and each variable, computes its unique reaching definition, and the summary associated set of reaching paths. This set of paths is abstracted using a *path expression*. If the code is not already under SSA, at a merge point of

<sup>1</sup> As opposed to the  $\psi$ -function described in Chapter ?? that would use a syntax such as  $a_3 = \phi((p \wedge \neg q)?a_1, (\neg p \wedge q)?a_2)$  instead.

the CFG, if its predecessor basic blocks are reached by different variables, a  $\phi$ -operator is inserted. The gates of each operand is set to the path expression of its corresponding incoming edge. If a unique variable reaches all the predecessor basic blocks, the corresponding path expressions are merged. Of course, classical path compression technique can be used to minimize the number of visited edges. One can observe the similarities with the  $\phi$ -operator placement algorithm described in Section ??.

~~A graph representation of our sample code in GSA form.~~ There also exists a relationship between the control dependencies and the gates: from a code already under strict and conventional SSA form, one can derive the gates of a  $\phi$  function from the control dependencies of its operands. This relationship is illustrated by Figure 1.4 in the simple case of a structured code.

These gating functions are important as the concept will form components of the Value State Dependence Graph later. GSA has seen a number of uses in the literature including analysis and transformations based on data flow. With the diversity of applications (see chapters ?? and ??), many variants of GSA have been proposed. Those variations concern the correct handling of loops in addition to the computation and representation of gates.

By using gating functions it becomes possible to construct IRs based solely on data dependencies. These IRs are sparse in nature compared to the CFG, making them good for analysis and transformation. This is also a more attractive proposition than generating and maintaining both a CFG and DFG, which can be complex and prone to human error. One approach has been to combine both of these into one representation, as is done in the PDG. Alternatively, we can utilize gating functions along with a data flow graph for an effective way of representing whole program information using data flow information.

### 1.5.1 Backwards symbolic analysis with GSA

GSA is useful for performing symbolic analysis. Traditionally, symbolic analysis is performed by forwards propagation of expressions through a program. However, complete forward substitution is expensive and can result in a large quantity of unused information and complicated expressions. ~~Tu and Padua showed how~~ Instead, backwards, demand-driven substitution can be performed using GSA which only substitutes *needed* information. Consider the following program, ~~taken from the authors:~~

```
R: JMAX = Expr
S: if(P) then J = JMAX - 1
   else J = JMAX
T: assert(J ≤ JMAX)
```

Fig. 1.5 A program on which to perform symbolic analysis.

If forwards substitution were to be used in order to determine whether the assertion is correct, then the symbolic value of  $J$  must be discovered, starting at the top of the program in statement  $R$ . Forward propagation through this program results in statement  $T$  being *assert*((if  $P$  then  $\text{Expr} - 1$  else  $\text{Expr}$ )  $\leq \text{Expr}$ ), thus the *assert* statement evaluates to true. In real, non-trivial programs, these expressions can get unnecessarily long and complicated.

Using GSA instead allows for backwards, demand-driven substitution. The program above has the following GSA form:

```

R: JMAX1 = Expr
S: if(P) then J1 = JMAX1 - 1
   else J2 = JMAX1
   J3 =  $\phi_{if}$ (P, J1, J2)
T: assert(J3 ≤ JMAX1)

```

**Fig. 1.6** Figure 1.5 in GSA form.

Using this backwards substitution technique, we start at statement  $T$ , and follow the SSA links of the variables from  $J_3$ . This allows for skipping of any intermediate statements that do not affect variables in  $T$ . Thus the substitution steps are:

$$\begin{aligned}
 J_3 &= \phi_{if}(P, J_1, J_2) \\
 &= \phi_{if}(P, JMAX_1 - 1, JMAX_1)
 \end{aligned}$$

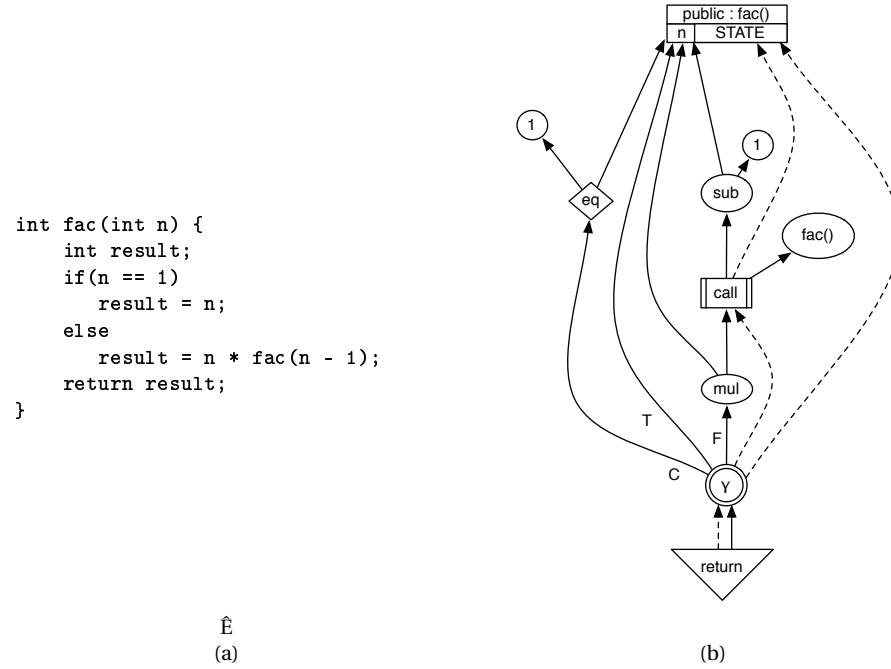
**Fig. 1.7** Substitution steps in backwards symbolic analysis.

The backwards substitution then stops because enough information has been found, avoiding the redundant substitution of  $JMAX_1$  by  $\text{Expr}$ . In non-trivial programs this can greatly reduce the number of redundant substitutions, making symbolic analysis significantly cheaper.

## 1.6 Value State Dependence Graph

The gating functions defined in the previous section were used in the development of a sparse data flow graph IR called the Value State Dependence Graph (VSDG). The VSDG is a directed graph consisting of operation nodes, loop and merge nodes together with value and state dependency edges. Cycles are permitted but must satisfy various restrictions. A VSDG represents a single procedure: this matches the classical CFG.

An example VSDG is shown in Figure 1.8. In (a) we have the original C source for a recursive factorial function. The corresponding VSDG (b) shows both value and state edges and a selection of nodes.



**Fig. 1.8** A recursive factorial function, whose VSDG illustrates the key graph components—value dependency edges (solid lines), state dependency edges (dashed lines), a `const` node, a `call` node, a  $\gamma$ -node, a conditional node, and the function entry and exit nodes.

### 1.6.1 Definition of the VSDG

A VSDG is a labelled directed graph  $G = (N, E_V, E_S, \ell, N_0, N_\infty)$  consisting of nodes  $N$  (with unique entry node  $N_0$  and exit node  $N_\infty$ ), value dependency edges  $E_V \subseteq N \times N$ , state dependency edges  $E_S \subseteq N \times N$ . The labelling function  $\ell$  associates each node with an operator.

The VSDG corresponds to a reducible program, e.g. there are no cycles in the VSDG except those mediated by  $\theta$  (loop) nodes.

Value dependency ( $E_V$ ) indicates the flow of values between nodes. State dependency ( $E_S$ ) represents two things; the first is essential sequential dependency required by the original program, e.g. a given `load` instruction may be required to follow a given `store` instruction without being re-ordered, and a `return` node

in general must wait for an earlier loop to terminate even though there might be no value-dependency between the loop and the `return` node. The second purpose is that state dependency edges can be added incrementally until the VSDG corresponds to a unique CFG. Such state dependency edges are called *serializing* edges.

The VSDG is implicitly represented in SSA form: a given operator node,  $n$ , will have zero or more  $E_V$ -consumers using its value. Note that, in implementation terms, a single register can hold the produced value for consumption at all consumers; it is therefore useful to talk about the idea of an output *port* for  $n$  being allocated a specific register,  $r$ , to abbreviate the idea of  $r$  being used for each edge  $(n_1, n_2)$  where  $n_2 \in \text{succ}(n_1)$ .

### 1.6.2 Nodes

There are four main classes of VSDG nodes: value nodes (representing pure arithmetic),  $\gamma$ -nodes (conditionals),  $\theta$ -nodes (loops), and state nodes (side-effects). The majority of nodes in a VSDG generate a value based on some computation (add, subtract, etc) applied to their dependent values (constant nodes, which have no dependent nodes, are a special case).

### 1.6.3 $\gamma$ -Nodes

The  $\gamma$ -node is similar to the  $\phi$ -gating function in being dependent on a control predicate, rather than the control-independent nature of SSA  $\phi$ -functions.

A  $\gamma$ -node  $\gamma(C, T, F)$  evaluates the condition dependency  $C$ , and returns the value of  $T$  if  $C$  is true, otherwise  $F$ .

We generally treat  $\gamma$ -nodes as single-valued nodes (contrast  $\theta$ -nodes, which are treated as tuples), with the effect that two separate  $\gamma$ -nodes with the same condition can be later combined into a tuple using a single test. Figure 1.9 illustrates two  $\gamma$ -nodes that can be combined in this way. Here, we use a pair of values (2-tuple) of values for the  $T$  and  $F$  ports. We also see how two syntactically different programs can map to the same structure in the VSDG.

### 1.6.4 $\theta$ -Nodes

The  $\theta$ -node models the iterative behaviour of loops, modelling loop state with the notion of an *internal value* which may be updated on each iteration of the loop. It has five specific ports which represent dependencies at various stages of computation.

A  $\theta$ -node  $\theta(C, I, R, L, X)$  sets its internal value to initial value  $I$  then, while condition value  $C$  holds true, sets  $L$  to the current internal value and updates the internal value with the repeat value  $R$ . When  $C$  evaluates to false computation ceases and the last internal value is returned through the  $X$  port.

A loop which updates  $k$  variables will have: a single condition port  $C$ , initial-value ports  $I_1, \dots, I_k$ , loop iteration ports  $L_1, \dots, L_k$ , loop return ports  $R_1, \dots, R_k$ , and loop exit ports  $X_1, \dots, X_k$ . The example in Figure 1.10 also shows a pair (2-tuple) of values being used for  $I, R, L, X$ , one for each loop-variant value.

The  $\theta$ -node directly implements pre-test loops (`while`, `for`); post-test loops (`do...while`, `repeat...until`) are synthesised from a pre-test loop preceded by a duplicate of the loop body. At first this may seem to cause unnecessary duplication of code, but it has two important benefits: (i) it exposes the first loop body iteration to optimization in post-test loops (cf. loop-peeling), and (ii) it normalizes all loops to one loop structure, which both reduces the cost of optimization, and increases the likelihood of two schematically-dissimilar loops being isomorphic in the VSDG.

### 1.6.5 State Nodes

Loads and stores compute a value and state. The `call` node takes both the name of the function to call and a list of arguments, and returns a list of results; it is treated as a state node as the function body may read or update state.

We maintain the simplicity of the VSDG by imposing the restriction that *all* functions have *one* return node (the exit node  $N_\infty$ ), which returns at least one result (which will be a state value in the case of `void` functions). To ensure that function calls and definitions are able to be allocated registers easily, we suppose that the number of arguments to, and results from, a function is smaller than the number of physical registers—further arguments can be passed via a stack as usual.

---


```

a) if (P)
    x = 2, y = 3;
else
    x = 4, y = 5;


b) if (P) x = 2; else x = 4;
    ...
    if (P) y = 3; else y = 5;

```

---



**Fig. 1.9** Two different code schemes (a) & (b) map to the same  $\gamma$ -node structure.



Note also that the VSDG neither forces loop invariant code into nor out-of loop bodies, but rather allows later phases to determine, by adding serializing edges, such placement of loop invariant nodes for later phases.

By representing a program as a VSDG, many optimisations become trivial. For example, consider dead node elimination (Figure 1.11). This combines both dead code elimination and unreachable code elimination. Dead code generates VSDG nodes for which there is no value or state dependency path from the `return` node, i.e., the result of the function does not in any way depend on the results of the dead nodes. Unreachable code generates VSDG nodes that are either dead, or become dead after some other optimisation. Thus, a *dead node* is a node that is not postdominated by the exit node  $N_\infty$ . To perform dead node elimination, only two passes are required over the VSDG resulting in linear runtime complexity: one pass to identify all of the live nodes, and a second pass to delete the unmarked (i.e., dead) nodes. It is safe because all nodes which are deleted are guaranteed never to be reachable from the `return` node.



**Input:** A VSDG  $G(N, E_V, E_S, N_\infty)$  with zero or more dead nodes.

**Output:** A VSDG with no dead nodes.

```

Procedure DNE( $G$ ) {
1:   WalkAndMark( $N_\infty, G$ );
2:   DeleteMarked( $G$ );
}
Procedure WalkAndMark( $n, G$ ) {
1:   if  $n$  is marked then finish;
2:   mark  $n$ ;
3:   foreach (node  $m \in N \wedge (n, m) \in (E_V \cup E_S)$ ) do
4:     WalkAndMark( $m$ );
}
Procedure DeleteMarked( $G$ ) {
1:   foreach (node  $n \in N$ ) do
2:     if  $n$  is unmarked then delete( $n$ );
}

```

Fig. 1.11 Dead node elimination on the VSDG.

## 1.7 Summary

A compiler's intermediate representation can be a graph, and many different graphs exist in the literature. We can represent the control flow of a program as a Control Flow Graph (CFG) [1], where straight-line instructions are contained within basic blocks and edges show where the flow of control may be transferred to once leaving that block. A CFG is traditionally used to convert a program to SSA form [7]. We can also represent programs as a type of Data Flow Graph (DFG) [9, 10], and SSA can be represented in this way as an SSA Graph [6]. An example was given that used the SSA Graph to detect a variety of induction variables in loops [31, 15]. It has also been used for performing instruction selection techniques [11, 22], operator strength reduction [6], rematerialization [5], and has been combined with an extended SSA language to aid compilation in a parallelizing compiler [25].

The Program Dependence Graph (PDG) as defined by Ferrante et al. [14] represents control and data dependencies in one graph. Their definition of control dependencies that turns out to be equivalent to post-dominance frontier lead to confusions at it uses a non standard definition of post-dominance. We choose to report the definition of Bilardi and Pingali [?]. Section ?? mentions possible abstractions to represent data dependencies for dynamically allocated objects. Among others, the book of Darte et al. [8] provides a good overview of such representations. The PDG has been used for program slicing [21], testing [3], and widely for parallelization [13, 12, 23, 4]. We showed an example of how the PDG directly exposes parallel code.

Gating functions can be used to create directly interpretable  $\phi$ -functions. These are used in Gated Single Assignment Form. An example was, Alpern et al. [2] presented a precursor of GSA for structured code, to detect equality of variables. This chapter adopts their notations, i.e. a  $\phi_{if}$  for a if-then-else construction,

a  $\phi_{entry}$  for the entry of a loop, and a  $\phi_{exit}$  for its exit. The original usage of GSA was by Ballance et al. [20] as an intermediate stage in the construction of the Program Dependence Web IR. Further GSA papers replaced  $\phi_{if}$  by  $\gamma$ ,  $\phi_{entry}$  by  $\mu$ , and  $\phi_{exit}$  by  $\eta$ . Havlak [16] presented an algorithm for construction of a simpler version of GSA—Thinned GSA—which is constructed from a CFG in SSA form. The construction technique sketched in this chapter is developed in more details in [26]. GSA has been used for a number of analyses and transformations based on data flow. The example given of how to perform backwards demand-driven symbolic analysis using GSA have been borrowed from [27]. If conversion (see Chapter ??), converts control dependencies into data dependencies. To avoid the potential loss of information related to the lowering of  $\phi$ -operators into conditional moves or select instructions, gating  $\psi$ -functions (see Chapter ??) can be used.

We then described the Value State Dependence Graph (VSDG) [18], which is an improvement on a previous, unmentioned graph, the Value Dependence Graph [30]. It uses the concept of gating functions, data dependencies and state to model a program. We gave an example of how to perform dead node elimination on the VSDG. Detailed semantics of the VSDG are available [18], as well as semantics of a related IR: the Gated Data Dependence Graph [29]. Further study has taken place on the problem of generating code from the VSDG [28, 19, 24], and it has also been used to perform a combined register allocation and code motion algorithm [17].

This citation is to make chapters without citations build without error. Please ignore it: [?].

---

## References

---

1. Frances E. Allen. Control flow analysis. *SIGPLAN Not.*, 5(7):1–19, 1970. 17
2. B. Alpern, M. N. Wegman, and F. K. Zadeck. Detecting equality of variables in programs. In *Proceedings of the 15th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '88, pages 1–11. ACM, 1988. 17
3. Samuel Bates and Susan Horwitz. Incremental program testing using program dependence graphs. In *POPL '93: Proceedings of the 20th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 384–396, New York, NY, USA, 1993. ACM. 17
4. W. Baxter and H. R. Bauer, III. The program dependence graph and vectorization. In *POPL '89: Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 1–11, New York, NY, USA, 1989. ACM. 17
5. Preston Briggs, Keith D. Cooper, and Linda Torczon. Rematerialization. In *PLDI '92: Proceedings of the ACM SIGPLAN 1992 conference on Programming language design and implementation*, pages 311–321, New York, NY, USA, 1992. ACM. 17
6. Keith D. Cooper, L. Taylor Simpson, and Christopher A. Vick. Operator strength reduction. *ACM Trans. Program. Lang. Syst.*, 23(5):603–625, 2001. 17
7. Ron Cytron, Jeanne Ferrante, Barry K. Rosen, Mark N. Wegman, and F. Kenneth Zadeck. Efficiently computing static single assignment form and the control dependence graph. *ACM Trans. Program. Lang. Syst.*, 13(4):451–490, 1991. 17
8. Alain Darte, Yves Robert, and Frederic Vivien. *Scheduling and Automatic Parallelization*. Birkhauser Boston, 1st edition, 2000. 17
9. J. B. Dennis. First version of a data flow procedure language. In *Programming Symposium, Proceedings Colloque sur la Programmation*, pages 362–376, London, UK, 1974. Springer-Verlag. 17
10. J. B. Dennis. Data flow supercomputers. *Computer*, 13(11):48–56, 1980. 17
11. Dietmar Ebner, Florian Brandner, Bernhard Scholz, Andreas Krall, Peter Wiedermann, and Albrecht Kadlec. Generalized instruction selection using SSA-graphs. In *LCTES '08: Proceedings of the 2008 ACM SIGPLAN-SIGBED conference on languages, compilers, and tools for embedded systems*, pages 31–40, New York, NY, USA, 2008. ACM. 17
12. J. Ferrante, M. Mace, and B. Simons. Generating sequential code from parallel code. In *ICS '88: Proceedings of the 2nd international conference on Supercomputing*, pages 582–592, New York, NY, USA, 1988. ACM. 17
13. Jeanne Ferrante and Mary Mace. On linearizing parallel code. In *POPL '85: Proceedings of the 12th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*, pages 179–190, New York, NY, USA, 1985. ACM. 17
14. Jeanne Ferrante, Karl J. Ottenstein, and Joe D. Warren. The program dependence graph and its use in optimization. *ACM Trans. Program. Lang. Syst.*, 9(3):319–349, 1987. 17
15. Michael P. Gerlek, Eric Stoltz, and Michael Wolfe. Beyond induction variables: detecting and classifying sequences using a demand-driven SSA form. *ACM Trans. Program. Lang. Syst.*, 17(1):85–122, 1995. 17
16. Paul Havlak. Construction of thinned gated single-assignment form. In *In Proc. 6rd Workshop on Programming Languages and Compilers for Parallel Computing*, pages 477–499. Springer Verlag, 1993. 18

17. Neil Johnson and Alan Mycroft. Combined code motion and register allocation using the value state dependence graph. In *In Proc. 12th International Conference on Compiler Construction (CCS03) (April 2003)*, pages 1–16, 2003. 18
18. Neil E. Johnson. Code size optimization for embedded processors. Technical Report UCAM-CL-TR-607, University of Cambridge, Computer Laboratory, November 2004. 18
19. Alan C. Lawrence. Optimizing compilation with the Value State Dependence Graph. Technical Report UCAM-CL-TR-705, University of Cambridge, Computer Laboratory, December 2007. 18
20. Karl J. Ottenstein, Robert A. Ballance, and Arthur B. MacCabe. The program dependence web: a representation supporting control-, data-, and demand-driven interpretation of imperative languages. In *PLDI '90: Proceedings of the ACM SIGPLAN 1990 conference on Programming language design and implementation*, pages 257–271, New York, NY, USA, 1990. ACM. 18
21. Karl J. Ottenstein and Linda M. Ottenstein. The program dependence graph in a software development environment. In *SDE 1: Proceedings of the first ACM SIGSOFT/SIGPLAN software engineering symposium on Practical software development environments*, pages 177–184, New York, NY, USA, 1984. ACM. 17
22. Stefan Schäfer and Bernhard Scholz. Optimal chain rule placement for instruction selection based on SSA graphs. In *SCOPES '07: Proceedings of the 10th international workshop on Software & compilers for embedded systems*, pages 91–100, New York, NY, USA, 2007. ACM. 17
23. B. Simons, D. Alpern, and J. Ferrante. A foundation for sequentializing parallel code. In *SPAA '90: Proceedings of the second annual ACM symposium on Parallel algorithms and architectures*, pages 350–359, New York, NY, USA, 1990. ACM. 17
24. James Stanier. Ph.D. thesis. Technical report, University of Sussex, School of Informatics, October 2011. 18
25. Eric Stoltz, Michael P. Gerlek, and Michael Wolfe. Extended SSA with factored use-def chains to support optimization and parallelism. In *In Proceedings of the 27th Annual Hawaii International Conference on System Sciences*, pages 43–52, 1993. 17
26. Peng Tu and David Padua. Efficient building and placing of gating functions. In *PLDI '95: Proceedings of the ACM SIGPLAN 1995 conference on Programming language design and implementation*, pages 47–55, New York, NY, USA, 1995. ACM. 18
27. Peng Tu and David Padua. Gated SSA-based demand-driven symbolic analysis for parallelizing compilers. In *Proceedings of the 9th international conference on Supercomputing, ICS '95*, pages 414–423, New York, NY, USA, 1995. ACM. 18
28. Eben Upton. Optimal sequentialization of gated data dependence graphs is NP-complete. In *PDPTA*, pages 1767–1770, 2003. 18
29. Eben Upton. Compiling with data dependence graphs, 2006. 18
30. Daniel Weise, Roger F. Crew, Michael Ernst, and Bjarne Steensgaard. Value dependence graphs: representation without taxation. In *POPL '94: Proceedings of the 21st ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages 297–310, New York, NY, USA, 1994. ACM. 18
31. Michael Wolfe. Beyond induction variables. *SIGPLAN Not.*, 27(7):162–174, 1992. 17